


CMSSW - ARM vs Intel

on power consumption

1. ARM vs Intel

- power consumed in event processing
- comparison

2. Preliminary RAPL measurements



4x Intel(R) Core(TM)2 Quad CPU Q9400 @ 2.66GHz

- the plots account for power consumption (W) during the Event processing stage.
- 1, 2, 3 and 4 processes take roughly the same time to process the events, as expected
- when there is a physical cpu overcommitment, the time difference is:

5 processes: 22.75% > 1 process


8 processes: 204% > 1 process

4x Intel(R) Core(TM)2 Quad CPU Q9400 @ 2.66GHz

no hyperthreading

speed step technology

http://ark.intel.com/products/49490/Intel-Atom-processor-D525-1M-Cache-1_80-GHz



2x Intel Atom CPU D525 @ 1.8GHz

- the plots account for power consumption (W) during the Event processing stage.
- 1 and 2 processes take roughly the same time to process the events, as expected
- when there is a physical cpu overcommitment, the time difference is:

3 processes: 29% > 1 process


4 processes: 58% > 1 process

- under cpu overcommit, the increase of time is relatively smaller than what happens in quad and ARM (maybe due to hyperthreading)

2x Intel Atom CPU D525 @ 1.8GHz

hyper-threading (2 processing threads per physical core)

http://ark.intel.com/products/49490/Intel-Atom-processor-D525-1M-Cache-1_80-GHz



4x ARMv7-A Cortex 2.50 DMIPS/MHz

- the plots account for power consumption (W) during the Event processing stage.
- 1, 2, 3 and 4 processes take roughly the same time to process the events, as expected
- when there is a physical cpu overcommitment, the time difference is:

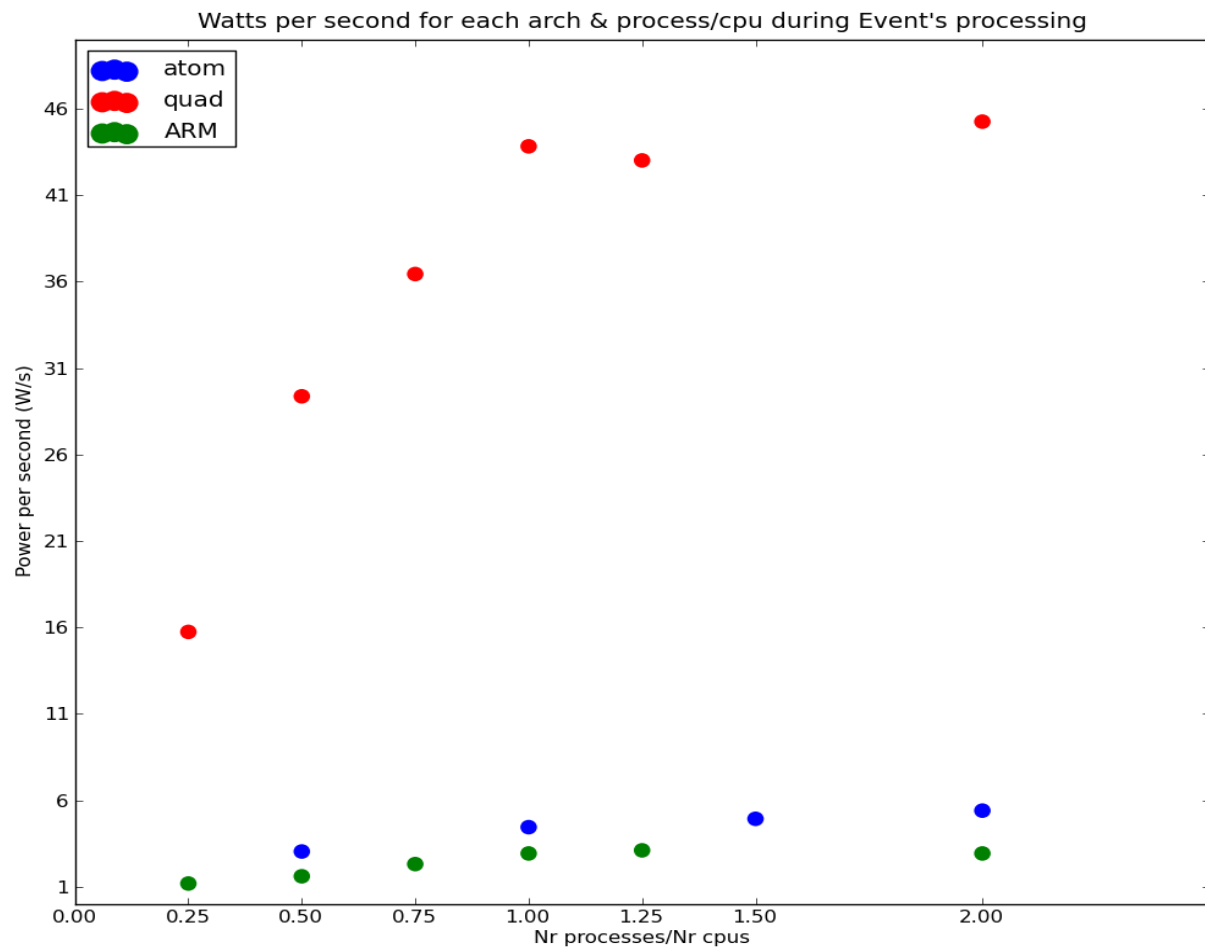
5 processes: 132% > 1 process

8 processes: 218% > 1 process

- plots don't have many samples, since the Event's processing stage is relatively small when compared with the startup phase.

4x ARMv7-A Cortex 2.50 DMIPS/MHz

<http://www.arm.com/products/processors/cortex-a/cortex-a9.php>



2. RAPL

running average power limit driver

- high precision power consumed in different domains: **socket**, **DRAM**, **processor**
- power consumption → better picture of how CMSSW behaves

machine specs caesrv1003.cern.ch

- Sandy Bridge machine, 4 CPUs total of 32 physical cores [2.7GHz, 20MB cache]
- Hyper-threading disabled
- 0.5TB of RAM

2. RAPL

bad performances if NUMA architecture is not taken into consideration at application level

→ What about energy efficiency ?

→ What about CMSSW ?

example

what happens if multithreading are scaled to more than 8 threads per socket ?

probably threads from same process running in different NUMA sockets will decrease significantly energy efficiency (due to latency, and others)

Some results


1. 16 processes

- i) bound to socket #2 and #3 (*using numactl -N ..*)
- ii) start processes normally

2. 32 processes


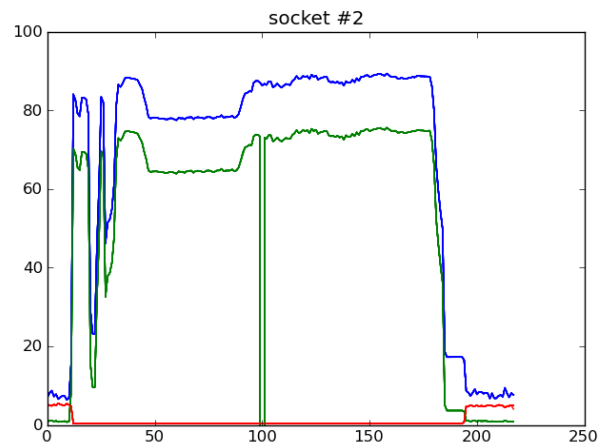


- i) bound to socket #2 and #3 (*using numactl -N ..*)
- ii) start processes normally

workflow: MinBias_8TeV_pythia8_cff_GEN_SIM.py, 10 Events, output off




16 processes bound to #2 and #3

- roughly 190s
- used DRAM is in #0 and #1
- kernel in #2 (*not confirmed*)


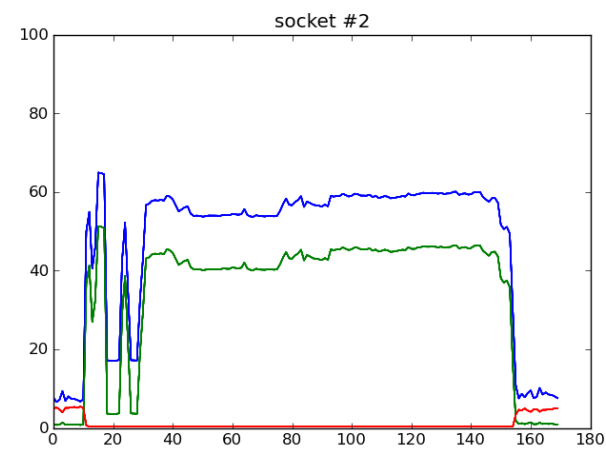




note: all processors are represented. though, they overlap per package




16 processes (no binding)

- roughly 150s
- used DRAM is in #0 and #1
- kernel in #0 (*not confirmed*)
- processes distributed evenly
- seems that when there is no explicit binding, the overall power efficiency is poorer (probably a too far fetched conclusion, will confirm it this week)


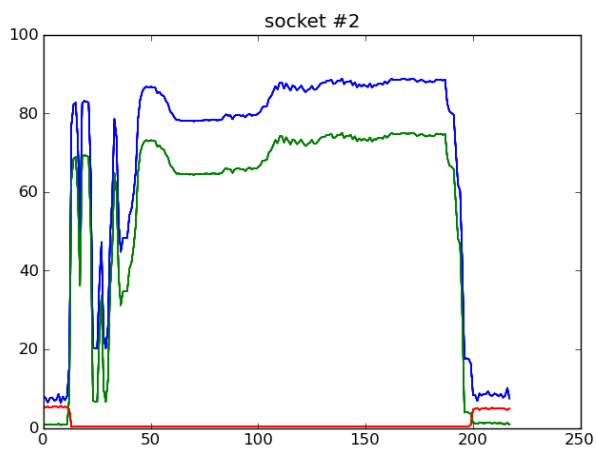




note: all processor are represented. though, they overlap per package




32 processes explicitly distributed with numactl (8 per socket)

- roughly 190s
- used DRAM is in #0 and #1
- not visually understandable where kernel is running



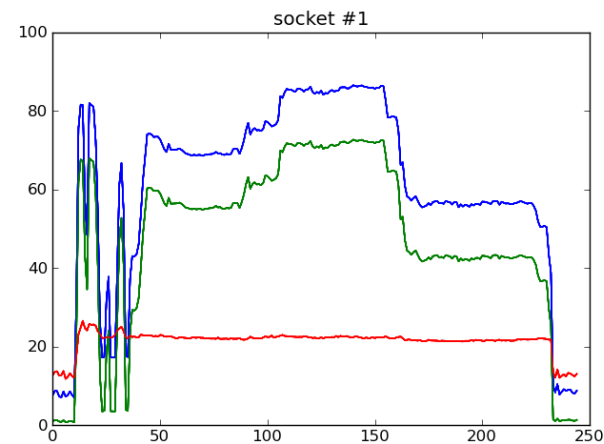



note: all processors are represented. though, they overlap per package



32 processes (no binding)

- roughly 240s
- used DRAM is in #0 and #1
- not visually understandable where kernel is running



note: all processors are represented. though, they overlap per package